

# Lexical Analysis

Dolors Sala

# Bibliography


These slides are taken from the following slide lectures

- Fredrik Kjolstad, Compilers CS143, Stanford University (Lectures 3 and 4)
- Henri Casanova, Machine learning and systems programming ICS312 (Module 9 Compiling), University of Hawaii (Lexing Lecture)
- Yannis Smaragdakis, Compilers K31, University of Athens, Lecture 3

# Lexical Analysis

- Identify the **words** in the input stream
- In linguistics the lexical analysis is to identify correct words
  - The correct words are listed at the dictionary
- The compiler with the lexical analysis does
  - **Identifies the words** (in the input code)
  - **Builds the table** (dictionary) of keywords and new identifiers (variable declarations, functions...)

# Lexical Analyzer: Definition



- Lexical analyzer translates the input code into a form more usable to the rest of the compiler
- It does two things:
  - Transforms the input source code into a sequence of substrings
  - Classifies them according to their role or syntactic category
- The input is the source code
- The output is a list of substrings called tokens
- Lexical analysis is also called lexing or scanning

# Token

- A **Token** is an indivisible **lexical unit** (a lexeme) classified according to syntactic categories
- **Lexeme:**
  - Keywords and special characters:
    - if, while, for,..
    - + - / : > < >> < ...
  - Names and numbers
- (Syntactic) **Category**: identifier, integer, floating-point number, operator, keyword...

# Lexical Analysis Terms

- A **Token** is an indivisible **lexical unit** (a lexeme) classified according to syntactic categories
- **Lexeme** is the string that represents an instance of a token
- **Pattern** is the description of the form that represents all possible lexemes a token can take

Example:

<i>Token</i>	<i>Pattern (informal description)</i>	<i>Sample lexemes</i>
<b>if</b>	characters <b>i</b> , <b>f</b>	<b>if</b>
<b>else</b>	characters <b>e</b> , <b>l</b> , <b>s</b> , <b>e</b>	<b>else</b>
<b>comparison</b>	< or > or <= or >= or == or !=	<b>&lt;=</b> , <b>!=</b>
<b>id</b>	letter followed by letters and digits	<b>pi</b> , <b>score</b> , <b>D2</b>
<b>number</b>	any numeric constant	<b>3.14159</b> , <b>0</b> , <b>6.02e23</b>
<b>literal</b>	anything but ", surrounded by ""s	<b>"core dumped"</b>

# Lexical Analysis Terminology Examples

<i>Token</i>	<i>Pattern (informal description)</i>	<i>Sample lexemes</i>
<b>if</b>	characters <b>i</b> , <b>f</b>	<b>if</b>
<b>else</b>	characters <b>e</b> , <b>l</b> , <b>s</b> , <b>e</b>	<b>else</b>
<b>comparison</b>	< or > or <= or >= or == or !=	<=, !=
<b>id</b>	letter followed by letters and digits	<b>pi</b> , <b>score</b> , <b>D2</b>
<b>number</b>	any numeric constant	<b>3.14159</b> , <b>0</b> , <b>6.02e23</b>
<b>literal</b>	anything but ", surrounded by ""s	" <b>core dumped</b> "

C example:

```
printf("Hello %s\n", name);
```

Tokens?

(printf, id)    ("Hello %s\n", literal)    (name, id)

How about ( , ; )?    They are also tokens, we'll see full specification

# The lookahead Problem

- The lexer has to process the source code **character by character** in order and identify the lexemes
- Just looking at the current character is not enough to determine if the lexeme is finished

Example: in c, = vs == or i vs int

In this case we need to check the next character to decide one or the other, this means a lookahead of 1, but... beginner vs begin?

- The **lookahead** is the amount of characters we need to look to correctly distinguish all possible lexemes of a language
  - Each language may have a different lookahead
  - Languages should require “small” lookahead

# Non-essential Characters

- The lexical analysis neglects non-essential characters
  - Spaces, tabs, linefeeds
  - Comments
- The lexer should allow arbitrary number of these non-essential characters in the input (source code)
- Input example:

```
main() {  
    int x = 0;  
}
```

Seen as a single string of characters: main() {\n\tint x = 0;\n}



# Lexical Analysis: Example (cont.)

- Input example:

```
main() {  
    int x = 0;  
}
```

The input is just a single string of characters:

```
main(){\n\tint x = 0;\n}
```

Tokenization:

(main, key)(“(”,op) (“)”,op) (“{”,op) (int, key)(x, id) (“=”,op) (“;”,op)  
 (“}”,op)

We'll see the exact specification of tokens format.

# Lexer Specification

- How to formalize the lexer's job to recognize the tokens of a specific language?
- We need a language: **the language of tokens**
- A language is specified with two components:
  - An **alphabet** denoted as  $\Sigma$ 
    - Example: ASCII
  - The **subset of all possible strings** over  $\Sigma$  (dictionary)
- Language of tokens -- specify which strings are tokens
  - It can be specified with **regular expressions**

# Examples of Languages

Alphabet = English characters      Alphabet = ASCII

Slight differences with Catalan/Spanish  
Big differences with Chinese, Arabic

Language = English sentences      Language = c programs

Not every string of English characters is an English sentence

Not every ASCII combination of characters is a correct c line of code

# Regular Expressions

- Languages are sets of strings
- Need a formal notation for specifying which ones
- The standard notation for regular languages is the regular expressions
- A **regular expression (RE)** is a string (in a meta-language) that describes a pattern (in the token language)
  - The regular expression has a notation/language so it is a language to describe a language (meta-language)

# Regular Expressions Specification

## Atomic regular expressions

- Single character: ‘c’ = {“c”}
- Epsilon :  $\epsilon = \{''\}$  The empty string (a string with 0 characters)

Compound regular expressions If A and B are regular expressions over an alphabet  $\Sigma$

- Union (+):  $A + B = \{s \mid s \in A \text{ or } s \in B\}$
- Concatenation :  $AB = \{ab \mid a \in A \text{ and } b \in B\}$
- Iteration (\*):  $A^* = \bigcup_{i \geq 0} A^i$  where  $A^i = AA\dots A$  i times

# More Regular Expressions Notation

Symbol	Meaning	Example
+	Union (one or another)	$a + b \rightarrow a \text{ or } b$
*	Zero or more times	$a^*$
$+ (a^+)$	One or more times	$a^+ \rightarrow aa^*$
?	Optional (zero or one time)	$b?a^* \rightarrow ba^* + a^*$
[]	Character class	[0-9] any digit [a-zA-Z_] any lower or upper case letter or underscore
?!	Negative lookahead	$(?!aa)a^*$ Any number of a's except aa
(?!pattern)		

## Useful notation for the P2 Lexer

# Examples of Regular Expressions (I)

- Regular expressions are common in our day-to-day life
- Phone numbers: (93)543-2333
  - $\Sigma = digits \cup \{-, '(', '\')\}$
  - exchange = digits<sup>2</sup>
  - area = digits<sup>3</sup>
  - phone = digits<sup>4</sup>
  - phone number = '(' + exchange + ')' + area + '-' + phone

Note the difference between a **meta-character** (a character of the language to define regular expressions) and a character of the alphabet: +34(93)543-2333

# Examples of Regular Expressions (II)

- Email Addresses: dolors@upf.edu
  - $\Sigma = \text{letter} \cup \{\cdot, '@'\}$  Defined in previous slides:  
letter = ‘A’ + ‘B’ + ... + ‘Z’ + ‘a’ + ‘b’ + ... + ‘z’
  - name = letter<sup>+</sup>
  - address = name + ‘@’ + name + ‘.’ + name
  
  
  
- Can it recognize the email: dolors.sala@upf.edu?
  - name = letter<sup>+</sup> + (‘.’ + letter<sup>+</sup>)\*
  
  
  
- Can it recognize the email: dolors2022@upf.edu?
  - $\Sigma = \text{letter} \cup \text{digit} \cup \{\cdot, '@'\}$
  - name = identifier<sup>+</sup> + ‘@’ + (‘.’ + identifier<sup>+</sup>)\*

Defined in previous slides:  
digit = [0-9]<sup>+</sup>  
identifier = letter (letter + digit)\*
  
  
  
- Can it recognize any email address?

# Notation: Tokens in regular expressions

- We describe tokens in regular expressions
- **Keyword:** “int” or “begin” or “else” or ...  
(Abbreviation: ‘int’ = ‘i’+ ‘n’+ ‘t’)
- **Integers:** a non-empty string of digits
  - $\Sigma = [0-9]$
  - digit = [0-9]
  - integer = digit<sup>+</sup>
- **Identifier:** strings of letters or digits, starting with a letter
  - $\Sigma = [a-zA-Z0-9]$
  - letter = [a-zA-Z]
  - identifier = letter (letter + digit)\*
- **Whitespace:** a non-empty sequence of blank spaces, newlines and tabs: (‘ ’ + ‘\n’ + ‘\t’)<sup>+</sup>

# Regular Expressions in SW


- Grep: Linux utility to identify a string in a file
- Perl Shell: program that process a text file applying a set of define string matching rules
- Text editors apply regular expressions for string replacements
- Many programs have the need to identify patterns and use regular expressions internally, including compilers
- **Not a unified syntax** of regular expressions: Each one has its own

# REs Implementation: Finite Automata

- How do we use the regular expressions (REs) to parse the input source code and generate the token stream?
- Regular expressions denote the (regular) language
- Regular languages are recognized by Finite Automata
- Therefore we implement the language defined by a regular expression implementing the corresponding automata

# Implementation of lexical Analysis

- Specify lexical structure in regular expressions
- Transform it to a finite automata
  - Deterministic finite automata DFA
  - Non-deterministic finite automata NFA
- Implementation of regular expressions



# Regular expressions “output”

- We can verify if a string belong to the language specified by the regular expression
  - True/false answer
- We need the tokens as the answer

# **Finite Automata**

## **(Implementation of regular expressions)**

A regular expression is implemented with an automaton

# Automatic Lexer Construction

- To convert a specification into code:
  1. Write down the RE for the input language
  2. Build a big NFA
  3. Build the DFA that simulates the NFA
  4. Systematically shrink the DFA
  5. Implement the DFA
- Lexer generators
  - Lex and flex work along these lines
  - Algorithms are well-known and well-understood
- In P2, we do some steps manually and create a generic engine to implement the automata as decided in each solution (one or many, DFA or NFA)
  - The automata are language specification provided as input

# Finite Automata

- Regular expressions = specification
- Finite automata = implementation
- A **finite automaton** is defined by
  - An input **alphabet**  $\Sigma$
  - A **set of states**  $S$
  - A **start state**  $n$  (only one state from the set  $S$ )
  - A set of **accepting states**  $F$  (a subset of  $S$ ,  $F \subseteq S$ )
  - A set of **transitions** between states: subset of  $S \times S$

# Transition notations

state: input → state

$s_1: a \rightarrow s_2$


state →<sup>input</sup> state

$s_1 \xrightarrow{a} s_2$


Both mean: if automaton is in state  $s_1$ , reading a character ‘a’ in the input takes the automaton in state  $s_2$  (we will use both indistinctively)

If when reaching the end of input, the automaton is in an accepting state => **accept** the input; otherwise **reject** it

# Finite Automata as Graphs




A **state** (state s)



The **start** state (state n)




An **accepting** state (state s)



A **transition** (from state s1 to state s2 with the a input)

# Automaton Example (I)


- What input does the following automaton accept?



It accepts the word “if”

# Automaton Example (II)


- What input does the following automaton accept?



It accepts the word 1

# Automaton Example (III)

- What input does the following automaton accept?



It accepts all strings that finish with 00

==> regular expression?  $(1+0)^*00$

# Lexical Analysis

Can we always represent a regular expression with an automaton?

YES!


If we write a piece of code that implements **arbitrary automata** we have a **generic piece** of code (an engine) that implements arbitrary regular expressions

This is a lexer!


# $\epsilon$ -Transitions

The epsilon ( $\epsilon$ ) transitions can move the automata from state to state **without consuming any input**  
 state:  $\epsilon \rightarrow$  state      state  $\rightarrow^{\epsilon}$  state


Defines a non-deterministic finite automata (NFA)  
 where there can be **multiple possible transitions from a given input symbol at a state**



# DFA vs NFA Example



DFA?



Any regular language can be defined by either NFA & DFA

# Deterministic and Non-deterministic FA

- A deterministic finite automata **DFA** is defined as an automaton that can only have one transition per input symbol per state
  - It can take **just one path** at any given time for a particular input
  - Accept the input if at the end of the input the automaton is at one accepting state
- Non-deterministic finite automata **NFA** is defined as an automaton that can have **zero ( $\epsilon$ ), one, or multiple transitions per input symbol per state**:
  - The NFA automaton have to **keep track of multiple paths simultaneously**
  - Accept the input if at the end of the input the automaton is **at least** at one accepting state

# DFA vs NFA Properties

- Both can describe the same set of languages (regular languages)
- DFAs are faster to execute (as there are no choices to consider)
- For some given languages the NFA can be simpler than DFA
- DFA can be exponentially larger than NFA

# Convert Regular Expressions to NFAs

- For each regular expression, define an NFA
- Thompson's construction (idea in 1968): systematically convert regular expressions (REs) into a finite state automata
- Key idea: Define a Finite Automata “pattern” for each RE operator
  - Start with atomic REs, build up a big NFA
- Notation: NFA for regular expression M



# Thompson's construction

- For each atomic regular expression, define an NFA
- Notation: NFA for regular expression M


- For  $\epsilon$




- For input a



- For AB



- For A+B



# Thompson's construction II


- For  $\epsilon$




- For input a




- For AB



- For A+B





- For  $A^*$





# Example of RegExp to NFA conversion I

- Regular Expression:  $(1+0)^*1$




# Example of RegExp to NFA conversion II

- Regular Expression:  $(1+0)^*1$




# Example of RegExp to NFA conversion III

- Regular Expression:  $(1+0)^*1$




# Example of RegExp to NFA conversion IV

- Regular Expression:  $(1+0)^*1$



# Example of RegExp to NFA conversion IV

- Regular Expression:  $(1+0)^*1$




- A human can define a simple NFA, however, this is a systematic approach for (automatic) implementation
  - The next step is to reduce it (systematically too)

## Step 2: From NFA to DFA


- Each state of DFA corresponds to a non-empty subset of states of NFA
- **Starting state** of the DFA is an state that contains all states of the NFA that are starting state and can be reached from the starting state with  $\epsilon$ -transitions
  - From A state can be reached BCDHI states → state 1 of DFA
- Add a **transition**  $S \xrightarrow{a} S'$  in a DFA state iff:
  - $S'$  is the set of NFA states reachable from any state in  $S$  after seeing the input  $a$ , considering  $\epsilon$ -transitions also
    - From state 1 with a 0 it goes to state 2 which corresponds in being in FGHIABCD
    - From state 1 with a 1 it goes to state 3 which corresponds being in EGHIJABCD (this is also a final state)
- **Final DFA states** are those DFA states which includes at least one NFA final state
  - State 3 is final because is the only one including state J

# Example NFA to DFA



# Implementation

- A DFA can be implemented by a 2D Table T
  - One dimension is the states
  - Other dimension is the input symbol
  - For every transition  $S_i \xrightarrow{a} S_j$  define  $T[i, a] = j$
- DFA execution:
  - If in state  $S_i$  and input  $a$ , read  $T[i, a] = k$  and go to state  $S_k$



	0	1
1	2	3
2	2	3
3	2	3

# Implementation

- NFA to DFA conversion is the heart of tools such as flex
- But DFAs can be huge
- In practice flex-like tools trade off speed for space in the choice of NFA and DFA implementation

# Implementing a Lexer

1. Define the regular expressions for each token category
2. Define an NFA for each RE
3. Convert the NFA (automatically) to a DFA
4. Write the code to implement the DFA
  1. Implement the transition table
5. Implement your lexer as a set of DFAs

All these steps have to be generic engines to apply the translation automatically to each input code.

These steps have been understood for decades and now there are the automatic lexer tools (lex, flex)

We don't implement all these automatic steps in practices

# Lexer Implementation

Two approaches:

- **By hand**: code by hand the lexer (in this course)
  - Start with the description for the lexemes of each token, and derive the automata of the language
  - Write code to identify each occurrence of each lexeme on the output to return the list of tokens identified
    - Write a generic engine that runs (NFA or DFA) automata
- **Automatically**: use a lexer generator (not in this course)
  - Specify the lexeme patterns to a lexical-analyzer generator
    - Write the engine that automatically identifies the automata based on the language description
  - Compile those patterns into code that works as lexical analyzer
    - Write the generic engine that runs the automata (as identified automatically)

# Summary

- The lexical analyzer identifies the **lexemes** (vocabulary) of the input and encodes them in **tokens** (lexeme, category)
- The identification of lexemes is done defining the **regular expression** that defines them
- Regular expressions are implemented by **automata**
  - NFA – non-deterministic finite automata
  - DFA – deterministic finite automata
- Lexical analysis process
  - Define: **alphabet, possible strings** (dictionary - table of tokens)
  - Define the Regular Expressions for each acceptable lexeme
  - Define the DFA that represent the REs (you may need to first define NFA and transform it to DFA)
  - Implement the DFA for each regular expression
  - Execute the DFA with an input to get the **tokens** or **errors**